

OTOT-A3

Name: Erin May Gunawan

Matric. No: A0200765L

GitHub Link: [A2-A3](#)

Demo Link: [A3](#)

Task A3.1

Install Metrics Server

1. Open your terminal and ensure the cluster is running and all the deployments, service, and ingress are setup

```
# Check if the kind nodes are running
docker ps -a

# Check the cluster info
kubectl cluster-info

# Check kubernetes resources
kubectl get <resource>
```

2. Install metrics server on cluster

```
kubectl apply -f https://github.com/kubernetes-sigs/metrics-
server/releases/latest/download/components.yaml
```

Edit the deployment configuration file for metric-server to add the flag `--kubelet-insecure-tls` in the `template.spec.containers[].args[]` field.

```
kubectl -n kube-system edit deploy/metrics-server
```

The final `yaml` should look something like this

```
# ...
template:
  # ...
  spec:
    containers:
      - args:
```

```
- --cert-dir=/tmp
- --secure-port=4443
- --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
- --kubelet-use-node-status-port
- --metric-resolution=15s
- --kubelet-insecure-tls

# ...
```

3. Check if the metric-server is running

```
# Check if deployed
kubectl -n kube-system get deploy/metrics-server

# sending the request to the API should return a result
# jq is optional and for formatting the output
kubectl get --raw "/apis/metrics.k8s.io/v1beta1/" | jq '.'
```

It should output something similar to this

```
{
  "kind": "APIResourceList",
  "apiVersion": "v1",
  "groupVersion": "metrics.k8s.io/v1beta1",
  "resources": [
    {
      "name": "nodes",
      "singularName": "",
      "namespaced": false,
      "kind": "NodeMetrics",
      "verbs": [
        "get",
        "list"
      ]
    },
    {
      "name": "pods",
      "singularName": "",
      "namespaced": true,
      "kind": "PodMetrics",
      "verbs": [
        "get",
        "list"
      ]
    }
  ]
}
```

Create Horizontal Pod Autoscaler

1. Create a manifest file for the Horizontal Pod Autoscaler, the content should be the following:

```
# backend-hpa.yaml

apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: backend
  namespace: default
spec:
  metrics:
    - resource:
        name: cpu
        target:
          averageUtilization: 50
          type: Utilization
        type: Resource
  minReplicas: 1
  maxReplicas: 10
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: backend
```

2. Deploy the HPA and verify that it is created correctly

```
# Deploy to cluster
kubectl apply -f k8s/manifests/backend-hpa.yaml

# Check if created
kubectl get hpa backend
```

It should output

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
backend	Deployment/backend	0%/50%	1	10	3	27s

The **TARGETS** is defined as **<current util> / <target util>** and should both be defined.

If **current util** exceeds the specified **target util**, the target deployment should scale accordingly.

```
# To get more details
kubectl describe hpa
```

It should output

```
HorizontalPodAutoscaler
Name:                                backend
Namespace:                          default
# ...
Reference:                           Deployment/backend
Metrics:                             ( current / target )
  resource cpu on pods  (as a percentage of request): 0% (0) / 50%
Min replicas:                        1
Max replicas:                        10
Deployment pods:                      3 current / 3 desired
```

The **Reference** and **Metrics** should be defined correctly.

You can also try opening up multiple tabs of the deployment via your browser and check if the deployment's **current utilization** changes and if it scales up.

```
# Watch for changes
kubectl get hpa backend -w
```

Or for stress-testing for multiple requests run the following in your terminal

```
# If using WSL2 to run, ensure the ip address is set to
# Ethernet adapter vEthernet (WSL)'s IP address
# check via running ifconfig on cmd

# Continously sends request, Ctrl+C to stop
while true; do curl <ip address of localhost>:<portNo>/app; done;
```

After some time running the command, the number of **REPLICAS** should increase. After stopping the command and waiting a couple of minutes, the number of **REPLICAS** should decrease.

Create Zone-Aware Deployment

1. Create the deployment

```
# backend-deployment-zone-aware.yaml

apiVersion: apps/v1
kind: Deployment
metadata:
  name: backend-zone-aware
  labels:
    app: backend-zone-aware
spec:
  replicas: 10
  selector:
```

```
matchLabels:
  app: backend-zone-aware
template:
  metadata:
    labels:
      app: backend-zone-aware
  spec:
    containers:
      - name: backend
        image: erinmayg/nginx-sample
        ports:
          - name: http
            containerPort: 80
        resources:
          limits:
            cpu: 40m
            memory: 100Mi
    topologySpreadConstraints:
      - maxSkew: 1
        topologyKey: topology.kubernetes.io/zone
        whenUnsatisfiable: DoNotSchedule
        labelSelector:
          matchLabels:
            app: backend-zone-aware
```

2. Deploy and verify

```
# Deploy
kubectl apply -f k8s/manifests/backend-deployment-zone-aware.yaml

# check if created
kubectl get po -lapp=backend-zone-aware -owide --sort-by='.spec.nodeName'
```